

```
In [1]: #!/ default_exp models.autoformer
```

Autoformer

The Autoformer model tackles the challenge of finding reliable dependencies on intricate temporal patterns of long-horizon forecasting.

The architecture has the following distinctive features:

- In-built progressive decomposition in trend and seasonal components based on a moving average filter.
- Auto-Correlation mechanism that discovers the period-based dependencies by calculating the autocorrelation and aggregating similar sub-series based on the periodicity.
- Classic encoder-decoder proposed by Vaswani et al. (2017) with a multi-head attention mechanism.

The Autoformer model utilizes a three-component approach to define its embedding:

- It employs encoded autoregressive features obtained from a convolution network.
- Absolute positional embeddings obtained from calendar features are utilized.

References

- [Wu, Haixu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. "Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting"](#)

Figure 1. Autoformer Architecture.

```
In [20]: #!/ export
import math
import numpy as np
from typing import Optional

import torch
import torch.nn as nn
import torch.nn.functional as F

from neuralforecast.common._modules import DataEmbedding
from neuralforecast.common._base_windows import BaseWindows

from neuralforecast.losses.pytorch import MAE
```

```
In [21]: #!/ hide
from fastcore.test import test_eq
from nbdev.showdoc import show_doc
```

1. Auxiliary Functions

```
In [22]: #!/ export
class AutoCorrelation(nn.Module):
    """
    AutoCorrelation Mechanism with the following two phases:
    (1) period-based dependencies discovery
    (2) time delay aggregation
    This block can replace the self-attention family mechanism seamlessly.
    """
    def __init__(self, mask_flag=True, factor=1, scale=None, attention_dropout=0.5):
        super(AutoCorrelation, self).__init__()
        self.factor = factor
        self.scale = scale
        self.mask_flag = mask_flag
        self.output_attention = output_attention
        self.dropout = nn.Dropout(attention_dropout)

    def time_delay_agg_training(self, values, corr):
        """
        SpeedUp version of Autocorrelation (a batch-normalization style design)
        This is for the training phase.
        """
        head = values.shape[1]
        channel = values.shape[2]
        length = values.shape[3]
        # find top k
        top_k = int(self.factor * math.log(length))
        mean_value = torch.mean(torch.mean(corr, dim=1), dim=1)
        index = torch.topk(torch.mean(mean_value, dim=0), top_k, dim=-1)[1]
        weights = torch.stack([mean_value[:, index[i]] for i in range(top_k)])
        # update corr
        tmp_corr = torch.softmax(weights, dim=-1)
        # aggregation
        tmp_values = values
        delays_agg = torch.zeros_like(values, dtype=torch.float, device=values.device)
        for i in range(top_k):
            pattern = torch.roll(tmp_values, -int(index[i]), -1)
            delays_agg = delays_agg + pattern * \
                (tmp_corr[:, i].unsqueeze(1).unsqueeze(1).unsqueeze(1))
        return delays_agg

    def time_delay_agg_inference(self, values, corr):
        """
        SpeedUp version of Autocorrelation (a batch-normalization style design)
        This is for the inference phase.
        """
        batch = values.shape[0]
        head = values.shape[1]
        channel = values.shape[2]
```

```

length = values.shape[3]
# index init
init_index = torch.arange(length, device=values.device).unsqueeze(0)
# find top k
top_k = int(self.factor * math.log(length))
mean_value = torch.mean(torch.mean(corr, dim=1), dim=1)
weights = torch.topk(mean_value, top_k, dim=-1)[0]
delay = torch.topk(mean_value, top_k, dim=-1)[1]
# update corr
tmp_corr = torch.softmax(weights, dim=-1)
# aggregation
tmp_values = values.repeat(1, 1, 1, 2)
delays_agg = torch.zeros_like(values, dtype=torch.float, device=values.device)
for i in range(top_k):
    tmp_delay = init_index + delay[:, i].unsqueeze(1).unsqueeze(1).unsqueeze(1)
    pattern = torch.gather(tmp_values, dim=-1, index=tmp_delay)
    delays_agg = delays_agg + pattern * \
        (tmp_corr[:, i].unsqueeze(1).unsqueeze(1).unsqueeze(1))
return delays_agg

def time_delay_agg_full(self, values, corr):
    """
    Standard version of Autocorrelation
    """
    batch = values.shape[0]
    head = values.shape[1]
    channel = values.shape[2]
    length = values.shape[3]
    # index init
    init_index = torch.arange(length, device=values.device).unsqueeze(0)
    # find top k
    top_k = int(self.factor * math.log(length))
    weights = torch.topk(corr, top_k, dim=-1)[0]
    delay = torch.topk(corr, top_k, dim=-1)[1]
    # update corr
    tmp_corr = torch.softmax(weights, dim=-1)
    # aggregation
    tmp_values = values.repeat(1, 1, 1, 2)
    delays_agg = torch.zeros_like(values, dtype=torch.float, device=values.device)
    for i in range(top_k):
        tmp_delay = init_index + delay[:, i].unsqueeze(-1)
        pattern = torch.gather(tmp_values, dim=-1, index=tmp_delay)
        delays_agg = delays_agg + pattern * (tmp_corr[:, i].unsqueeze(-1))
    return delays_agg

def forward(self, queries, keys, values, attn_mask):
    B, L, H, E = queries.shape
    _, S, _, D = values.shape
    if L > S:
        zeros = torch.zeros_like(queries[:, :(L - S), :], dtype=torch.float)
        values = torch.cat([values, zeros], dim=1)
        keys = torch.cat([keys, zeros], dim=1)
    else:
        values = values[:, :L, :, :]
        keys = keys[:, :L, :, :]

```

```

        # period-based dependencies
        q_fft = torch.fft.rfft(queries.permute(0, 2, 3, 1).contiguous(), dim=-1)
        k_fft = torch.fft.rfft(keys.permute(0, 2, 3, 1).contiguous(), dim=-1)
        res = q_fft * torch.conj(k_fft)
        corr = torch.fft.irfft(res, dim=-1)

        # time delay agg
        if self.training:
            V = self.time_delay_agg_training(values.permute(0, 2, 3, 1).contiguous(), corr)
        else:
            V = self.time_delay_agg_inference(values.permute(0, 2, 3, 1).contiguous(), corr)

        if self.output_attention:
            return (V.contiguous(), corr.permute(0, 3, 1, 2))
        else:
            return (V.contiguous(), None)

class AutoCorrelationLayer(nn.Module):
    def __init__(self, correlation, hidden_size, n_head, d_keys=None,
                  d_values=None):
        super(AutoCorrelationLayer, self).__init__()

        d_keys = d_keys or (hidden_size // n_head)
        d_values = d_values or (hidden_size // n_head)

        self.inner_correlation = correlation
        self.query_projection = nn.Linear(hidden_size, d_keys * n_head)
        self.key_projection = nn.Linear(hidden_size, d_keys * n_head)
        self.value_projection = nn.Linear(hidden_size, d_values * n_head)
        self.out_projection = nn.Linear(d_values * n_head, hidden_size)
        self.n_head = n_head

    def forward(self, queries, keys, values, attn_mask):
        B, L, _ = queries.shape
        _, S, _ = keys.shape
        H = self.n_head

        queries = self.query_projection(queries).view(B, L, H, -1)
        keys = self.key_projection(keys).view(B, S, H, -1)
        values = self.value_projection(values).view(B, S, H, -1)

        out, attn = self.inner_correlation(
            queries,
            keys,
            values,
            attn_mask
        )
        out = out.view(B, L, -1)

        return self.out_projection(out), attn

class LayerNorm(nn.Module):
    """
    Special designed layernorm for the seasonal part

```

```

"""
def __init__(self, channels):
    super(LayerNorm, self).__init__()
    self.layernorm = nn.LayerNorm(channels)

def forward(self, x):
    x_hat = self.layernorm(x)
    bias = torch.mean(x_hat, dim=1).unsqueeze(1).repeat(1, x.shape[1], 1)
    return x_hat - bias

class MovingAvg(nn.Module):
    """
    Moving average block to highlight the trend of time series
    """
    def __init__(self, kernel_size, stride):
        super(MovingAvg, self).__init__()
        self.kernel_size = kernel_size
        self.avg = nn.AvgPool1d(kernel_size=kernel_size, stride=stride, padding=0)

    def forward(self, x):
        # padding on the both ends of time series
        front = x[:, 0:1, :].repeat(1, (self.kernel_size - 1) // 2, 1)
        end = x[:, -1:, :].repeat(1, (self.kernel_size - 1) // 2, 1)
        x = torch.cat([front, x, end], dim=1)
        x = self.avg(x.permute(0, 2, 1))
        x = x.permute(0, 2, 1)
        return x

class SeriesDecomp(nn.Module):
    """
    Series decomposition block
    """
    def __init__(self, kernel_size):
        super(SeriesDecomp, self).__init__()
        self.MovingAvg = MovingAvg(kernel_size, stride=1)

    def forward(self, x):
        moving_mean = self.MovingAvg(x)
        res = x - moving_mean
        return res, moving_mean

class EncoderLayer(nn.Module):
    """
    Autoformer encoder layer with the progressive decomposition architecture
    """
    def __init__(self, attention, hidden_size, conv_hidden_size=None, MovingAvg=None):
        super(EncoderLayer, self).__init__()
        conv_hidden_size = conv_hidden_size or 4 * hidden_size
        self.attention = attention
        self.conv1 = nn.Conv1d(in_channels=hidden_size, out_channels=conv_hidden_size, kernel_size=3, padding=1)
        self.conv2 = nn.Conv1d(in_channels=conv_hidden_size, out_channels=hidden_size, kernel_size=3, padding=1)
        self.decomp1 = SeriesDecomp(MovingAvg)
        self.decomp2 = SeriesDecomp(MovingAvg)

```

```

        self.dropout = nn.Dropout(dropout)
        self.activation = F.relu if activation == "relu" else F.gelu

    def forward(self, x, attn_mask=None):
        new_x, attn = self.attention(
            x, x, x,
            attn_mask=attn_mask
        )
        x = x + self.dropout(new_x)
        x, _ = self.decomp1(x)
        y = x
        y = self.dropout(self.activation(self.conv1(y.transpose(-1, 1))))
        y = self.dropout(self.conv2(y).transpose(-1, 1))
        res, _ = self.decomp2(x + y)
        return res, attn

class Encoder(nn.Module):
    """
    Autoformer encoder
    """
    def __init__(self, attn_layers, conv_layers=None, norm_layer=None):
        super(Encoder, self).__init__()
        self.attn_layers = nn.ModuleList(attn_layers)
        self.conv_layers = nn.ModuleList(conv_layers) if conv_layers is not None else None
        self.norm = norm_layer

    def forward(self, x, attn_mask=None):
        attns = []
        if self.conv_layers is not None:
            for attn_layer, conv_layer in zip(self.attn_layers, self.conv_layers):
                x, attn = attn_layer(x, attn_mask=attn_mask)
                x = conv_layer(x)
                attns.append(attn)
            x, attn = self.attn_layers[-1](x)
            attns.append(attn)
        else:
            for attn_layer in self.attn_layers:
                x, attn = attn_layer(x, attn_mask=attn_mask)
                attns.append(attn)

        if self.norm is not None:
            x = self.norm(x)

        return x, attns

class DecoderLayer(nn.Module):
    """
    Autoformer decoder layer with the progressive decomposition architecture
    """
    def __init__(self, self_attention, cross_attention, hidden_size, c_out,
                 MovingAvg=25, dropout=0.1, activation="relu"):
        super(DecoderLayer, self).__init__()
        conv_hidden_size = conv_hidden_size or 4 * hidden_size
        self.self_attention = self_attention

```

```

self.cross_attention = cross_attention
self.conv1 = nn.Conv1d(in_channels=hidden_size, out_channels=conv_hi
self.conv2 = nn.Conv1d(in_channels=conv_hidden_size, out_channels=hi
self.decomp1 = SeriesDecomp(MovingAvg)
self.decomp2 = SeriesDecomp(MovingAvg)
self.decomp3 = SeriesDecomp(MovingAvg)
self.dropout = nn.Dropout(dropout)
self.projection = nn.Conv1d(in_channels=hidden_size, out_channels=c
padding_mode='circular', bias=False)
self.activation = F.relu if activation == "relu" else F.gelu

def forward(self, x, cross, x_mask=None, cross_mask=None):
    x = x + self.dropout(self.self_attention(
        x, x, x,
        attn_mask=x_mask
    )[0])
    x, trend1 = self.decomp1(x)
    x = x + self.dropout(self.cross_attention(
        x, cross, cross,
        attn_mask=cross_mask
    )[0])
    x, trend2 = self.decomp2(x)
    y = x
    y = self.dropout(self.activation(self.conv1(y.transpose(-1, 1))))
    y = self.dropout(self.conv2(y).transpose(-1, 1))
    x, trend3 = self.decomp3(x + y)

    residual_trend = trend1 + trend2 + trend3
    residual_trend = self.projection(residual_trend.permute(0, 2, 1)).tr
    return x, residual_trend

```

```

class Decoder(nn.Module):
    """
    Autoformer decoder
    """
    def __init__(self, layers, norm_layer=None, projection=None):
        super(Decoder, self).__init__()
        self.layers = nn.ModuleList(layers)
        self.norm = norm_layer
        self.projection = projection

    def forward(self, x, cross, x_mask=None, cross_mask=None, trend=None):
        for layer in self.layers:
            x, residual_trend = layer(x, cross, x_mask=x_mask, cross_mask=cr
            trend = trend + residual_trend

        if self.norm is not None:
            x = self.norm(x)

        if self.projection is not None:
            x = self.projection(x)
        return x, trend

```

3 Autoformer

In [23]: `#!/ export`

```
class Autoformer(BaseWindows):  
    """ Autoformer
```

The Autoformer model tackles the challenge of finding reliable dependence

The architecture has the following distinctive features:

- In-built progressive decomposition in trend and seasonal components b
- Auto-Correlation mechanism that discovers the period-based dependence calculating the autocorrelation and aggregating similar sub-series based
- Classic encoder-decoder proposed by Vaswani et al. (2017) with a multi

The Autoformer model utilizes a three-component approach to define its e

- It employs encoded autoregressive features obtained from a convolution
- Absolute positional embeddings obtained from calendar features are uti

Parameters:

``h`: int, forecast horizon.
`

``input_size`: int, maximum sequence length for truncated train backpropa`

``futr_exog_list`: str list, future exogenous columns.
`

``hist_exog_list`: str list, historic exogenous columns.
`

``stat_exog_list`: str list, static exogenous columns.
`

``exclude_insample_y`: bool=False, the model skips the autoregressive fea`

``decoder_input_size_multiplier`: float = 0.5, .
`

``hidden_size`: int=128, units of embeddings and encoders.
`

``n_head`: int=4, controls number of multi-head's attention.
`

``dropout`: float (0, 1), dropout throughout Autoformer architecture.
`

``factor`: int=3, Probsparse attention factor.
`

``conv_hidden_size`: int=32, channels of the convolutional encoder.
`

``activation`: str='GELU', activation from ['ReLU', 'Softplus', 'Tanh`

``encoder_layers`: int=2, number of layers for the TCN encoder.
`

``decoder_layers`: int=1, number of layers for the MLP decoder.
`

``distil`: bool = True, wether the Autoformer decoder uses bottlenecks.
`

``loss`: PyTorch module, instantiated train loss class from [losses colle`

``max_steps`: int=1000, maximum number of training steps.
`

``learning_rate`: float=1e-3, Learning rate between (0, 1).
`

``num_lr_decays`: int=-1, Number of learning rate decays, evenly distribu`

``early_stop_patience_steps`: int=-1, Number of validation iterations bef`

``val_check_steps`: int=100, Number of training steps between every valid`

``batch_size`: int=32, number of different series in each batch.
`

``valid_batch_size`: int=None, number of different series in each validat`

``windows_batch_size`: int=1024, number of windows to sample in each trai`

``inference_windows_batch_size`: int=1024, number of windows to sample in`

``start_padding_enabled`: bool=False, if True, the model will pad the tim`

``scaler_type`: str='robust', type of scaler for temporal inputs normaliz`

``random_seed`: int=1, random_seed for pytorch initializer and numpy gene`

``num_workers_loader`: int=os.cpu_count(), workers to be used by `TimeSer`

``drop_last_loader`: bool=False, if True `TimeSeriesDataLoader` drops las`

``alias`: str, optional, Custom name of the model.
`

`**trainer_kwargs`: int, keyword trainer arguments inherited from [PyTo`

References

- [Wu, Haixu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. "Autoform

"""


```
SAMPLING_TYPE = 'windows'
```

```
def __init__(self,
              h: int,
              input_size: int,
              stat_exog_list = None,
              hist_exog_list = None,
              futr_exog_list = None,
              exclude_insample_y = False,
              decoder_input_size_multiplier: float = 0.5,
              hidden_size: int = 128,
              dropout: float = 0.05,
              factor: int = 3,
              n_head: int = 4,
              conv_hidden_size: int = 32,
              activation: str = 'gelu',
              encoder_layers: int = 2,
              decoder_layers: int = 1,
              MovingAvg_window: int = 25,
              loss = MAE(),
              valid_loss = None,
              max_steps: int = 5000,
              learning_rate: float = 1e-4,
              num_lr_decays: int = -1,
              early_stop_patience_steps: int = -1,
              val_check_steps: int = 100,
              batch_size: int = 32,
              valid_batch_size: Optional[int] = None,
              windows_batch_size = 1024,
              inference_windows_batch_size = 1024,
              start_padding_enabled = False,
              step_size: int = 1,
              scaler_type: str = 'identity',
              random_seed: int = 1,
              num_workers_loader: int = 0,
              drop_last_loader: bool = False,
              **trainer_kwargs):
    super(Autoformer, self).__init__(h=h,
                                     input_size=input_size,
                                     hist_exog_list=hist_exog_list,
                                     stat_exog_list=stat_exog_list,
                                     futr_exog_list = futr_exog_list,
                                     exclude_insample_y = exclude_insample_y,
                                     loss=loss,
                                     valid_loss=valid_loss,
                                     max_steps=max_steps,
                                     learning_rate=learning_rate,
                                     num_lr_decays=num_lr_decays,
                                     early_stop_patience_steps=early_stop,
                                     val_check_steps=val_check_steps,
                                     batch_size=batch_size,
                                     windows_batch_size=windows_batch_size,
                                     valid_batch_size=valid_batch_size,
                                     inference_windows_batch_size=inferenc
                                     start_padding_enabled = start_padding
                                     step_size=step_size,
```

```

        scaler_type=scaler_type,
        num_workers_loader=num_workers_loader,
        drop_last_loader=drop_last_loader,
        random_seed=random_seed,
        **trainer_kwargs)

    # Architecture
    self.futr_input_size = len(self.futr_exog_list)
    self.hist_input_size = len(self.hist_exog_list)
    self.stat_input_size = len(self.stat_exog_list)

    if self.stat_input_size > 0:
        raise Exception('Autoformer does not support static variables yet')

    if self.hist_input_size > 0:
        raise Exception('Autoformer does not support historical variables yet')

    self.label_len = int(np.ceil(input_size * decoder_input_size_multiplier))
    if (self.label_len >= input_size) or (self.label_len <= 0):
        raise Exception(f'Check decoder_input_size_multiplier={decoder_input_size_multiplier}')

    if activation not in ['relu', 'gelu']:
        raise Exception(f'Check activation={activation}')

    self.c_out = self.loss.outputsize_multiplier
    self.output_attention = False
    self.enc_in = 1
    self.dec_in = 1

    # Decomposition
    self.decomp = SeriesDecomp(MovingAvg_window)

    # Embedding
    self.enc_embedding = DataEmbedding(c_in=self.enc_in,
                                       exog_input_size=self.hist_input_size,
                                       hidden_size=hidden_size,
                                       pos_embedding=False,
                                       dropout=dropout)
    self.dec_embedding = DataEmbedding(self.dec_in,
                                       exog_input_size=self.hist_input_size,
                                       hidden_size=hidden_size,
                                       pos_embedding=False,
                                       dropout=dropout)

    # Encoder
    self.encoder = Encoder(
        [
            EncoderLayer(
                AutoCorrelationLayer(
                    AutoCorrelation(False, factor,
                                    attention_dropout=dropout,
                                    output_attention=self.output_attention,
                                    hidden_size, n_head),
                    hidden_size=hidden_size,
                    conv_hidden_size=conv_hidden_size,
                    MovingAvg=MovingAvg_window,

```

```

        dropout=dropout,
        activation=activation
    ) for l in range(encoder_layers)
],
    norm_layer=LayerNorm(hidden_size)
)
# Decoder
self.decoder = Decoder(
    [
        DecoderLayer(
            AutoCorrelationLayer(
                AutoCorrelation(True, factor, attention_dropout=dropo
                hidden_size, n_head),
            AutoCorrelationLayer(
                AutoCorrelation(False, factor, attention_dropout=dro
                hidden_size, n_head),
            hidden_size=hidden_size,
            c_out=self.c_out,
            conv_hidden_size=conv_hidden_size,
            MovingAvg=MovingAvg_window,
            dropout=dropout,
            activation=activation,
        )
        for l in range(decoder_layers)
    ],
    norm_layer=LayerNorm(hidden_size),
    projection=nn.Linear(hidden_size, self.c_out, bias=True)
)

def forward(self, windows_batch):
    # Parse windows_batch
    insample_y = windows_batch['insample_y']
    #insample_mask = windows_batch['insample_mask']
    #hist_exog = windows_batch['hist_exog']
    #stat_exog = windows_batch['stat_exog']
    futr_exog = windows_batch['futr_exog']

    # Parse inputs
    insample_y = insample_y.unsqueeze(-1) # [Ws,L,1]
    if self.futr_input_size > 0:
        x_mark_enc = futr_exog[:, :, self.input_size, :]
        x_mark_dec = futr_exog[:, -(self.label_len+self.h), :, :]
    else:
        x_mark_enc = None
        x_mark_dec = None

    x_dec = torch.zeros(size=(len(insample_y), self.h, 1)).to(insample_y.device)
    x_dec = torch.cat([insample_y[:, -self.label_len:, :, :], x_dec], dim=1)

    # decomp init
    mean = torch.mean(insample_y, dim=1).unsqueeze(1).repeat(1, self.h, 1)
    zeros = torch.zeros([x_dec.shape[0], self.h, x_dec.shape[2]], device=x_dec.device)
    seasonal_init, trend_init = self.decomp(insample_y)
    # decoder input
    trend_init = torch.cat([trend_init[:, -self.label_len:, :], mean], dim=1)
    seasonal_init = torch.cat([seasonal_init[:, -self.label_len:, :], zeros], dim=1)

```

```

# enc
enc_out = self.enc_embedding(insample_y, x_mark_enc)
enc_out, attns = self.encoder(enc_out, attn_mask=None)
# dec
dec_out = self.dec_embedding(seasonal_init, x_mark_dec)
seasonal_part, trend_part = self.decoder(dec_out, enc_out, x_mask=None,
                                         trend=trend_init)

# final
dec_out = trend_part + seasonal_part

forecast = self.loss.domain_map(dec_out[:, -self.h:])
return forecast

```

In [24]: `show_doc(Autoformer)`

Autoformer

```

Autoformer (h:int, input_size:int, stat_exog_list=None,
            hist_exog_list=None, futr_exog_list=None,
            exclude_insample_y=False,
            decoder_input_size_multiplier:float=0.5, hidd
en_size:int=128,
            dropout:float=0.05, factor:int=3, n_head:int=
4,
            conv_hidden_size:int=32, activation:str='gel
u',
            encoder_layers:int=2, decoder_layers:int=1,
            MovingAvg_window:int=25, loss=MAE(), valid_lo
ss=None,
            max_steps:int=5000, learning_rate:float=0.000
1,
            num_lr_decays:int=-1, early_stop_patience_ste
ps:int=-1,
            val_check_steps:int=100, batch_size:int=32,
            valid_batch_size:Optional[int]=None, windows_
batch_size=1024,
            inference_windows_batch_size=1024,
            start_padding_enabled=False, step_size:int=1,
            scaler_type:str='identity', random_seed:int=
1,
            num_workers_loader:int=0, drop_last_loader:bo
ol=False,
            **trainer_kwargs)

```

Autoformer

The Autoformer model tackles the challenge of finding reliable dependencies on intricate temporal patterns of long-horizon forecasting.

The architecture has the following distinctive features:

- In-built progressive decomposition in trend and seasonal components based on a moving average filter.
- Auto-Correlation mechanism that discovers the period-based dependencies by calculating the autocorrelation and aggregating similar sub-series based on the periodicity.

- Classic encoder-decoder proposed by Vaswani et al. (2017) with a multi-head attention mechanism.

The Autoformer model utilizes a three-component approach to define its embedding:

- It employs encoded autoregressive features obtained from a convolution network.
- Absolute positional embeddings obtained from calendar features are utilized.

Parameters:

`h` : int, forecast horizon.

`input_size` : int, maximum sequence length for truncated train backpropagation. Default -1 uses all history.

`futr_exog_list` : str list, future exogenous columns.

`hist_exog_list` : str list, historic exogenous columns.

`stat_exog_list` : str list, static exogenous columns.

`exclude_insample_y` : bool=False, the model skips the autoregressive features $y[t-input_size:t]$ if True.

`decoder_input_size_multiplier` : float = 0.5, .

`hidden_size` : int=128, units of embeddings and encoders.

`n_head` : int=4, controls number of multi-head's attention.

`dropout` : float (0, 1), dropout throughout Autoformer architecture.

`factor` : int=3, Probsparse attention factor.

`conv_hidden_size` : int=32, channels of the convolutional encoder.

`activation` : str= GELU , activation from ['ReLU', 'Softplus', 'Tanh', 'SELU', 'LeakyReLU', 'PReLU', 'Sigmoid', 'GELU'].

`encoder_layers` : int=2, number of layers for the TCN encoder.

`decoder_layers` : int=1, number of layers for the MLP decoder.

`distil` : bool = True, whether the Autoformer decoder uses bottlenecks.

`loss` : PyTorch module, instantiated train loss class from [losses collection](#).

`max_steps` : int=1000, maximum number of training steps.

`learning_rate` : float=1e-3, Learning rate between (0, 1).

`num_lr_decays` : int=-1, Number of learning rate decays, evenly distributed across `max_steps`.

`early_stop_patience_steps` : int=-1, Number of validation iterations before early stopping.

`val_check_steps` : int=100, Number of training steps between every validation loss check.

`batch_size` : int=32, number of different series in each batch.

`valid_batch_size` : int=None, number of different series in each validation and test batch, if None uses `batch_size`.

`windows_batch_size` : int=1024, number of windows to sample in each training

batch, default uses all.

`inference_windows_batch_size` : int=1024, number of windows to sample in each inference batch.

`start_padding_enabled` : bool=False, if True, the model will pad the time series with zeros at the beginning, by input size.

`scaler_type` : str='robust', type of scaler for temporal inputs normalization see [temporal scalers](#).

`random_seed` : int=1, random_seed for pytorch initializer and numpy generators.

`num_workers_loader` : int=os.cpu_count(), workers to be used by

`TimeSeriesDataLoader` .

`drop_last_loader` : bool=False, if True `TimeSeriesDataLoader` drops last non-full batch.

`alias` : str, optional, Custom name of the model.

`**trainer_kwargs` : int, keyword trainer arguments inherited from [PyTorch Lightning's trainer](#).

References

- [Wu, Haixu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. "Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting"](<https://proceedings.neurips.cc/paper/2021/hash/bcc0d400288793e8bdcd7c19a8ac0c2b-Abstract.html>)


```
In [6]: show_doc(Autoformer.fit, name='Autoformer.fit')
```

Out[6]:

Autoformer.fit

```
Autoformer.fit (dataset, val_size=0, test_size=0, random_
seed=None)
```

Fit.

The `fit` method, optimizes the neural network's weights using the initialization parameters (`learning_rate` , `windows_batch_size` , ...) and the `loss` function as defined during the initialization. Within `fit` we use a PyTorch Lightning `Trainer` that inherits the initialization's `self.trainer_kwargs` , to customize its inputs, see [PL's trainer arguments](#).

The method is designed to be compatible with SKLearn-like classes and in particular to be compatible with the StatsForecast library.

By default the `model` is not saving training checkpoints to protect disk memory, to get them change `enable_checkpointing=True` in `__init__`.

Parameters:

`dataset` : NeuralForecast's `TimeSeriesDataset` , see [documentation](#).

`val_size` : int, validation size for temporal cross-validation.

`random_seed` : int=None, random_seed for pytorch initializer and numpy generators, overwrites model.`init`'s.

`test_size` : int, test size for temporal cross-validation.

```
In [8]: show_doc(Autoformer.predict, name='Autoformer.predict')
```


Out[8]:

Autoformer.predict

```
Autoformer.predict (dataset, test_size=None, step_size=1,
                    random_seed=None, **data_module_kwarg
s)
```

Predict.

Neural network prediction with PL's `Trainer` execution of `predict_step`.

Parameters:

`dataset` : NeuralForecast's `TimeSeriesDataset`, see [documentation](#).
`test_size` : int=None, test size for temporal cross-validation.
`step_size` : int=1, Step size between each window.
`random_seed` : int=None, random_seed for pytorch initializer and numpy generators, overwrites model.`init`'s.
`**data_module_kwargs` : PL's `TimeSeriesDataModule` args, see [documentation](#).

Usage Example

```
In [9]: #!/ eval: false
import numpy as np
import pandas as pd
import pytorch_lightning as pl
import matplotlib.pyplot as plt

from neuralforecast import NeuralForecast
from neuralforecast.models import MLP
from neuralforecast.losses.pytorch import MQLoss, DistributionLoss
from neuralforecast.tsdataset import TimeSeriesDataset
from neuralforecast.utils import AirPassengers, AirPassengersPanel, AirPassengersPanel

AirPassengersPanel, calendar_cols = augment_calendar_df(df=AirPassengersPanel

Y_train_df = AirPassengersPanel[AirPassengersPanel.ds<AirPassengersPanel['ds
Y_test_df = AirPassengersPanel[AirPassengersPanel.ds>=AirPassengersPanel['ds

model = Autoformer(h=12,
                  input_size=24,
                  hidden_size = 16,
                  conv_hidden_size = 32,
                  n_head=2,
                  loss=MAE(),
                  futr_exog_list=calendar_cols,
                  scaler_type='robust',
                  learning_rate=1e-3,
                  max_steps=300,
                  val_check_steps=50,
```

```

        early_stop_patience_steps=2)

nf = NeuralForecast(
    models=[model],
    freq='M'
)
nf.fit(df=Y_train_df, static_df=AirPassengersStatic, val_size=12)
forecasts = nf.predict(futr_df=Y_test_df)

Y_hat_df = forecasts.reset_index(drop=False).drop(columns=['unique_id', 'ds'])
plot_df = pd.concat([Y_test_df, Y_hat_df], axis=1)
plot_df = pd.concat([Y_train_df, plot_df])

if model.loss.is_distribution_output:
    plot_df = plot_df[plot_df.unique_id=='Airline1'].drop('unique_id', axis=1)
    plt.plot(plot_df['ds'], plot_df['y'], c='black', label='True')
    plt.plot(plot_df['ds'], plot_df['Autoformer-median'], c='blue', label='m')
    plt.fill_between(x=plot_df['ds'][-12:],
                     y1=plot_df['Autoformer-lo-90'][-12:].values,
                     y2=plot_df['Autoformer-hi-90'][-12:].values,
                     alpha=0.4, label='level 90')

    plt.grid()
    plt.legend()
    plt.plot()
else:
    plot_df = plot_df[plot_df.unique_id=='Airline1'].drop('unique_id', axis=1)
    plt.plot(plot_df['ds'], plot_df['y'], c='black', label='True')
    plt.plot(plot_df['ds'], plot_df['Autoformer'], c='blue', label='Forecast')
    plt.legend()
    plt.grid()

```

Seed set to 1

2023-11-02 00:51:42.424894: I tensorflow/core/util/port.cc:111] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.

2023-11-02 00:51:42.515256: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not find cuda drivers on your machine, GPU will not be used.

2023-11-02 00:51:42.837663: E tensorflow/compiler/xla/stream_executor/cuda/cuda_dnn.cc:9342] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered

2023-11-02 00:51:42.837728: E tensorflow/compiler/xla/stream_executor/cuda/cuda_fft.cc:609] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered

2023-11-02 00:51:42.839420: E tensorflow/compiler/xla/stream_executor/cuda/cuda_blas.cc:1518] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered

2023-11-02 00:51:42.970876: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

2023-11-02 00:51:44.751923: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT


Sanity Checking: |

10/250000...

```

Training: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Predicting: |
| 0/? [00:00...

```



```

In [18]: y_true = Y_test_df.y.values
          y_hat = Y_hat_df['Autoformer'].values

          from neuralforecast.losses.numpy import mae, mse

          print('MAE: ', mae(y_hat, y_true))
          print('MSE: ', mse(y_hat, y_true))

```

```

MAE: 15.130399068196615
MSE: 349.1350725169759

```

Autoformer Implementation

Exchange rate

```
In [19]: import pandas as pd
from neuralforecast import NeuralForecast

Y_df = pd.read_csv("raw_data/df_Exchange.csv")

Y_df['ds'] = pd.to_datetime(Y_df['ds'])

# For this excercise we are going to take 20% of the DataSet
n_time = len(Y_df.ds.unique())
val_size = int(.1 * n_time)
test_size = int(.2 * n_time)

Y_df.groupby('unique_id').head(2)
```

```
Out[19]:
```

	unique_id	ds	y
0	0	1990-01-01	0.606785
1	0	1990-01-02	0.570900
7588	1	1990-01-01	-0.361671
7589	1	1990-01-02	-0.367639
15176	2	1990-01-01	0.735367
15177	2	1990-01-02	0.729629
22764	3	1990-01-01	-1.164373
22765	3	1990-01-02	-1.170907
30352	4	1990-01-01	2.851890
30353	4	1990-01-02	2.851890
37940	5	1990-01-01	-1.861369
37941	5	1990-01-02	-1.838665
45528	6	1990-01-01	-1.820047
45529	6	1990-01-02	-1.847258
53116	OT	1990-01-01	-0.124081
53117	OT	1990-01-02	-0.113588

```
In [25]: horizon = 96

model = Autoformer(h=horizon,
                  input_size = horizon,
                  max_steps=100,
                  val_check_steps=10,
                  batch_size = 8,
                  hidden_size = 32,
                  windows_batch_size = 256,
```

```

        early_stop_patience_steps=2)

nf = NeuralForecast(
    models=[model],
    freq='D'
)

Y_hat_df = nf.cross_validation(df=Y_df,
                              val_size=val_size,
                              test_size=test_size,
                              n_windows=None)

```

Seed set to 1

```

Sanity Checking: |
| 0/? [00:00...
Training: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Predicting: |
| 0/? [00:00...

```

In [27]: `Y_hat_df.head()`

```

Out[27]:

```

	unique_id	ds	cutoff	Autoformer	y
0	0	2006-08-16	2006-08-15	3.044964	2.948076
1	0	2006-08-17	2006-08-15	3.030476	3.049320
2	0	2006-08-18	2006-08-15	3.028900	3.064168
3	0	2006-08-19	2006-08-15	3.050043	3.005783
4	0	2006-08-20	2006-08-15	3.117382	3.010031

In [28]: `Y_hat_df.to_csv('results/Exchange_rate/autoformer.csv')`

In [26]: `from neuralforecast.losses.numpy import mae, mse`

```

print('MAE: ', mae(Y_hat_df['y'], Y_hat_df['Autoformer']))
print('MSE: ', mse(Y_hat_df['y'], Y_hat_df['Autoformer']))

```

```

MAE: 0.29645489153468074
MSE: 0.1861597390966

```

Weather

```
In [11]: import pandas as pd
         from neuralforecast import NeuralForecast

         Y_df = pd.read_csv("raw_data/df_Weather.csv")

         Y_df['ds'] = pd.to_datetime(Y_df['ds'])

         # For this exercise we are going to take 20% of the DataSet
         n_time = len(Y_df.ds.unique())
         val_size = int(.1 * n_time)
         test_size = int(.2 * n_time)

         Y_df.groupby('unique_id').head(2)
```

Out[11]:

	unique_id	ds	y
0	H2OC (mmol/mol)	2020-01-01 00:10:00	-0.999107
1	H2OC (mmol/mol)	2020-01-01 00:20:00	-1.008072
52695	OT	2020-01-01 00:10:00	0.044395
52696	OT	2020-01-01 00:20:00	0.044134
105390	PAR ($\mu\text{mol/m}^2/\text{s}$)	2020-01-01 00:10:00	-0.679493
105391	PAR ($\mu\text{mol/m}^2/\text{s}$)	2020-01-01 00:20:00	-0.679493
158085	SWDR (W/m^2)	2020-01-01 00:10:00	-0.672767
158086	SWDR (W/m^2)	2020-01-01 00:20:00	-0.672767
210780	T (degC)	2020-01-01 00:10:00	-1.459980
210781	T (degC)	2020-01-01 00:20:00	-1.454798
263475	Tdew (degC)	2020-01-01 00:10:00	-1.052596
263476	Tdew (degC)	2020-01-01 00:20:00	-1.069612
316170	Tlog (degC)	2020-01-01 00:10:00	-1.424132
316171	Tlog (degC)	2020-01-01 00:20:00	-1.416612
368865	Tpot (K)	2020-01-01 00:10:00	-1.607935
368866	Tpot (K)	2020-01-01 00:20:00	-1.602882
421560	VPact (mbar)	2020-01-01 00:10:00	-0.979132
421561	VPact (mbar)	2020-01-01 00:20:00	-0.990506
474255	VPdef (mbar)	2020-01-01 00:10:00	-0.838497
474256	VPdef (mbar)	2020-01-01 00:20:00	-0.828332
526950	VPmax (mbar)	2020-01-01 00:10:00	-1.141181
526951	VPmax (mbar)	2020-01-01 00:20:00	-1.138714
579645	max. PAR ($\mu\text{mol/m}^2/\text{s}$)	2020-01-01 00:10:00	-0.588296
579646	max. PAR ($\mu\text{mol/m}^2/\text{s}$)	2020-01-01 00:20:00	-0.588296
632340	max. wv (m/s)	2020-01-01 00:10:00	-0.832381
632341	max. wv (m/s)	2020-01-01 00:20:00	-1.125140
685035	p (mbar)	2020-01-01 00:10:00	2.114257
685036	p (mbar)	2020-01-01 00:20:00	2.099194
737730	rain (mm)	2020-01-01 00:10:00	-0.093506
737731	rain (mm)	2020-01-01 00:20:00	-0.093506
790425	raining (s)	2020-01-01 00:10:00	-0.221050
790426	raining (s)	2020-01-01 00:20:00	-0.221050
843120	rh (%)	2020-01-01 00:10:00	0.990128

	unique_id	ds	y
843121	rh (%)	2020-01-01 00:20:00	0.942141
895815	rho (g/m**3)	2020-01-01 00:10:00	1.940406
895816	rho (g/m**3)	2020-01-01 00:20:00	1.932788
948510	sh (g/kg)	2020-01-01 00:10:00	-0.998513
948511	sh (g/kg)	2020-01-01 00:20:00	-1.009228
1001205	wd (deg)	2020-01-01 00:10:00	0.555571
1001206	wd (deg)	2020-01-01 00:20:00	0.354339
1053900	wv (m/s)	2020-01-01 00:10:00	-0.017801
1053901	wv (m/s)	2020-01-01 00:20:00	-0.029125

In [8]: horizon = 96

```

model = Autoformer(h=horizon,
                    input_size = horizon,
                    max_steps=100,
                    val_check_steps=10,
                    batch_size = 21,
                    hidden_size = 16,
                    windows_batch_size = 256,
                    early_stop_patience_steps=2)

nf = NeuralForecast(
    models=[model],
    freq='10min'
)

Y_hat_df = nf.cross_validation(df=Y_df,
                               val_size=val_size,
                               test_size=test_size,
                               n_windows=None)

```



```
Seed set to 1
2023-11-02 05:44:45.812388: I tensorflow/core/util/port.cc:111] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2023-11-02 05:44:45.813943: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not find cuda drivers on your machine, GPU will not be used.
2023-11-02 05:44:45.836010: E tensorflow/compiler/xla/stream_executor/cuda/cuda_dnn.cc:9342] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
2023-11-02 05:44:45.836041: E tensorflow/compiler/xla/stream_executor/cuda/cuda_fft.cc:609] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
2023-11-02 05:44:45.836057: E tensorflow/compiler/xla/stream_executor/cuda/cuda_blas.cc:1518] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
2023-11-02 05:44:45.841303: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2023-11-02 05:44:46.432899: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT
```

```
Sanity Checking: |
```

```
| 0/? [00:00...
```

```
Training: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Validation: |
```

```
| 0/? [00:00...
```

```
Predicting: |
```

```
| 0/? [00:00...
```

```
In [10]: Y_hat_df.to_csv('results/Weather/Autoformer.csv')
```

```
In [11]: from neuralforecast.losses.numpy import mae, mse
```

```
print('MAE: ', mae(Y_hat_df['y'], Y_hat_df['Autoformer']))
print('MSE: ', mse(Y_hat_df['y'], Y_hat_df['Autoformer']))
```

MAE: 0.28974621539184015
MSE: 0.2305957080949505

Ettm2

```
In [7]: import pandas as pd
        from neuralforecast import NeuralForecast

        Y_df = pd.read_csv("raw_data/df_Ettm2.csv")

        Y_df['ds'] = pd.to_datetime(Y_df['ds'])

        # For this exercise we are going to take 20% of the DataSet
        n_time = len(Y_df.ds.unique())
        val_size = int(.2 * n_time)
        test_size = int(.2 * n_time)

        Y_df.groupby('unique_id').head(2)
```

Out[7]:

	unique_id	ds	y
0	HUFL 2016-07-01 00:00:00	-0.041413	
1	HUFL 2016-07-01 00:15:00	-0.185467	
57600	HULL 2016-07-01 00:00:00	0.040104	
57601	HULL 2016-07-01 00:15:00	-0.214450	
115200	LUFL 2016-07-01 00:00:00	0.695804	
115201	LUFL 2016-07-01 00:15:00	0.434685	
172800	LULL 2016-07-01 00:00:00	0.434430	
172801	LULL 2016-07-01 00:15:00	0.428168	
230400	MUFL 2016-07-01 00:00:00	-0.599211	
230401	MUFL 2016-07-01 00:15:00	-0.658068	
288000	MULL 2016-07-01 00:00:00	-0.393536	
288001	MULL 2016-07-01 00:15:00	-0.659338	
345600	OT 2016-07-01 00:00:00	1.018032	
345601	OT 2016-07-01 00:15:00	0.980124	

```
In [8]: horizon = 96

        model = Autoformer(h=horizon,
                           input_size = horizon,
                           max_steps=100,
                           val_check_steps=10,
                           batch_size = 7,
```

```

        hidden_size = 16,
        windows_batch_size = 256,
        early_stop_patience_steps=2)

nf = NeuralForecast(
    models=[model],
    freq='15min'
)

Y_hat_df = nf.cross_validation(df=Y_df,
                               val_size=val_size,
                               test_size=test_size,
                               n_windows=None)

```

```

Seed set to 1
2023-11-02 06:05:55.346787: I tensorflow/core/util/port.cc:111] oneDNN custom
operations are on. You may see slightly different numerical results due to
floating-point round-off errors from different computation orders. To turn t
hem off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2023-11-02 06:05:55.416697: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could n
ot find cuda drivers on your machine, GPU will not be used.
2023-11-02 06:05:55.854010: E tensorflow/compiler/xla/stream_executor/cuda/c
uda_dnn.cc:9342] Unable to register cuDNN factory: Attempting to register fa
ctory for plugin cuDNN when one has already been registered
2023-11-02 06:05:55.854073: E tensorflow/compiler/xla/stream_executor/cuda/c
uda_fft.cc:609] Unable to register cuFFT factory: Attempting to register fac
tory for plugin cuFFT when one has already been registered
2023-11-02 06:05:55.856131: E tensorflow/compiler/xla/stream_executor/cuda/c
uda_blas.cc:1518] Unable to register cuBLAS factory: Attempting to register
factory for plugin cuBLAS when one has already been registered
2023-11-02 06:05:56.039684: I tensorflow/core/platform/cpu_feature_guard.cc:
182] This TensorFlow binary is optimized to use available CPU instructions i
n performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operation
s, rebuild TensorFlow with the appropriate compiler flags.
2023-11-02 06:05:58.652714: W tensorflow/compiler/tf2tensorrt/utils/py_util
s.cc:38] TF-TRT Warning: Could not find TensorRT

```

```

Sanity Checking: |
| 0/? [00:00...
Training: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...

```

```
Validation: |  
| 0/? [00:00...  
Validation: |  
| 0/? [00:00...  
Validation: |  
| 0/? [00:00...  
Predicting: |  
| 0/? [00:00...
```

```
In [9]: Y_hat_df.to_csv('results/Ettm2/Autoformer.csv')
```

```
In [11]: from neuralforecast.losses.numpy import mae, mse  
  
print('MAE: ', mae(Y_hat_df['y'], Y_hat_df['Autoformer']))  
print('MSE: ', mse(Y_hat_df['y'], Y_hat_df['Autoformer']))
```

```
MAE: 0.3353140537163157  
MSE: 0.2553693649575985
```

ILI

```
In [12]: import pandas as pd  
from neuralforecast import NeuralForecast  
  
Y_df = pd.read_csv("raw_data/df_ILI.csv")  
  
Y_df['ds'] = pd.to_datetime(Y_df['ds'])  
  
# For this exercise we are going to take 20% of the DataSet  
n_time = len(Y_df.ds.unique())  
val_size = int(.1 * n_time)  
test_size = int(.2 * n_time)  
  
Y_df.groupby('unique_id').head(2)
```

Out[12]:

	unique_id	ds	y
0	% WEIGHTED ILI	2002-01-01	-0.421499
1	% WEIGHTED ILI	2002-01-08	-0.331239
966	%UNWEIGHTED ILI	2002-01-01	-0.472442
967	%UNWEIGHTED ILI	2002-01-08	-0.429154
1932	AGE 0-4	2002-01-01	-0.981641
1933	AGE 0-4	2002-01-08	-0.934213
2898	AGE 5-24	2002-01-01	-0.692621
2899	AGE 5-24	2002-01-08	-0.676837
3864	ILITOTAL	2002-01-01	-0.819695
3865	ILITOTAL	2002-01-08	-0.796703
4830	NUM. OF PROVIDERS	2002-01-01	-1.151274
4831	NUM. OF PROVIDERS	2002-01-08	-1.088458
5796	OT	2002-01-01	-1.385709
5797	OT	2002-01-08	-1.342939

In [16]:

```
horizon = 24

model = Autoformer(h=horizon,
                   input_size = horizon,
                   max_steps=100,
                   val_check_steps=10,
                   batch_size = 7,
                   hidden_size = 16,
                   windows_batch_size = 256,
                   early_stop_patience_steps=2)

nf = NeuralForecast(
    models=[model],
    freq='W'
)

Y_hat_df = nf.cross_validation(df=Y_df,
                              val_size=val_size,
                              test_size=test_size,
                              n_windows=None)
```

Seed set to 1

Sanity Checking: |

| 0/? [00:00...

Training: |

| 0/? [00:00...

Validation: |

| 0/? [00:00...

Validation: |

Loading [MathJax]/extensions/Safe.js)0...

```

Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Validation: |
| 0/? [00:00...
Predicting: |
| 0/? [00:00...

```

In [17]: Y_hat_df

```

Out[17]:

```

	unique_id	ds	cutoff	Autoformer	y
0	% WEIGHTED ILI	2016-10-23	2016-10-16	0.751263	NaN
1	% WEIGHTED ILI	2016-10-30	2016-10-16	0.743696	NaN
2	% WEIGHTED ILI	2016-11-06	2016-10-16	0.570308	NaN
3	% WEIGHTED ILI	2016-11-13	2016-10-16	0.405683	NaN
4	% WEIGHTED ILI	2016-11-20	2016-10-16	0.290363	NaN
...
28555	OT	2020-05-31	2020-01-12	4.320040	NaN
28556	OT	2020-06-07	2020-01-12	4.685857	NaN
28557	OT	2020-06-14	2020-01-12	4.701055	NaN
28558	OT	2020-06-21	2020-01-12	4.397973	NaN
28559	OT	2020-06-28	2020-01-12	2.905941	NaN

28560 rows × 5 columns

In []: Y_hat_df.to_csv('results/ILI/Autoformer.csv')

```

In [18]: from neuralforecast.losses.numpy import mae, mse

print('MAE: ', mae(Y_hat_df['y'], Y_hat_df['Autoformer']))
print('MSE: ', mse(Y_hat_df['y'], Y_hat_df['Autoformer']))

```

```

MAE:  nan
MSE:  nan

```

In []: